

# Network — IPv4 Tunnel over TCP/IPv6

SAAOUDI Mohamed

January 19, 2026

## Contents

<b>1 Organization, Demonstration</b>	<b>2</b>
1.1 Recommended Directory Structure . . . . .	2
<b>2 Network Configuration (Part 1)</b>	<b>2</b>
2.1 Topology and Addressing Plan . . . . .	2
2.2 VM2 Failure and Impact . . . . .	2
<b>3 TUN Virtual Interface (Part 2)</b>	<b>2</b>
3.1 Interface Creation (2.1) . . . . .	2
3.2 Configuration of tun0 (2.2) . . . . .	2
3.2.1 Script <code>configure-tun.sh</code> (2.2.1) . . . . .	2
3.2.2 Routing: VM1 or VM1-6? (2.2.2) . . . . .	4
3.2.3 Local ping tun0 (2.2.3) . . . . .	4
3.2.4 Non-local ping via tun0 (2.2.4) and explanation (2.2.5) . . . . .	4
3.3 Packet Retrieval (2.3) . . . . .	5
3.3.1 Option IFF_NO_PI (2.3.4) . . . . .	5
<b>4 Simple IPv4 Tunnel over TCP/IPv6 (Part 3)</b>	<b>5</b>
4.1 Incoming Traffic Redirection (3.1) . . . . .	5
4.2 Final Integration (3.3) . . . . .	5
4.3 VM1-6 <-> VM3-6 Tunnel: Diagrams (3.4) . . . . .	5
<b>5 Functional Validation (Part 4)</b>	<b>5</b>
5.1 Configuration (4.1) . . . . .	5
5.2 Layer 3 (4.2) . . . . .	10
5.3 Layer 4 (4.3) . . . . .	10
5.4 Layer 4: Bandwidth (4.4) . . . . .	11
<b>6 Conclusion</b>	<b>11</b>

# 1 Organization, Demonstration

## 1.1 Recommended Directory Structure

- `src/` : libraries (`iftun`, `extremite`, `traitement`) + `tunnel46d`
- `script/` : network configuration + demonstration scenarios
- `captures/` : screenshots + `.pcap` files (tshark/Wireshark)
- `fig/` : report figures
- `rapport/` : `rapport.tex` + compiled PDF

# 2 Network Configuration (Part 1)

## 2.1 Topology and Addressing Plan

The addressing plan used is the one from the instructions: LAN1..LAN4 in IPv4 (/28) and two IPv6 LANs `fc00:1234:1::/64` and `fc00:1234:2::/64`.<sup>1</sup>

Typical addresses (excerpt):

- VM1 : 172.16.2.151/28 (LAN3)
- VM3 : 172.16.2.183/28 (LAN4)
- VM1-6 : 172.16.2.156/28 (LAN3) and IPv6 `fc00:1234:1::16`
- VM3-6 : 172.16.2.186/28 (LAN4) and IPv6 `fc00:1234:2::36`

Figure 1 summarizes the topology and the key addresses used.

## 2.2 VM2 Failure and Impact

Since VM2 is stopped, VM1 can no longer reach VM3 directly. The IPv4-over-TCP/IPv6 tunnel via VM1-6 and VM3-6 restores connectivity.<sup>2</sup>

The initial ping failure (VM2 down) is visible in Figure 2.

# 3 TUN Virtual Interface (Part 2)

## 3.1 Interface Creation (2.1)

The TUN interface is created via a program (library `iftun`) based on the `tunalloc.c` code (non-persistent interface: it disappears when the creating process stops).<sup>3</sup>

The creation of the `tun0` interface is visible in Figure 3.

## 3.2 Configuration of tun0 (2.2)

### 3.2.1 Script `configure-tun.sh` (2.2.1)

The instructions require configuring `tun0` with the address 172.16.2.1 and an appropriate mask, via a script.<sup>4</sup>

**Example script (to be adapted to your VM):**

<sup>1</sup>Addressing table "1.1 Topology and Addressing".

<sup>2</sup>Context "1.2 A great misfortune!".

<sup>3</sup>Instructions "2.1 Creation of the interface".

<sup>4</sup>Point "2.2.1".

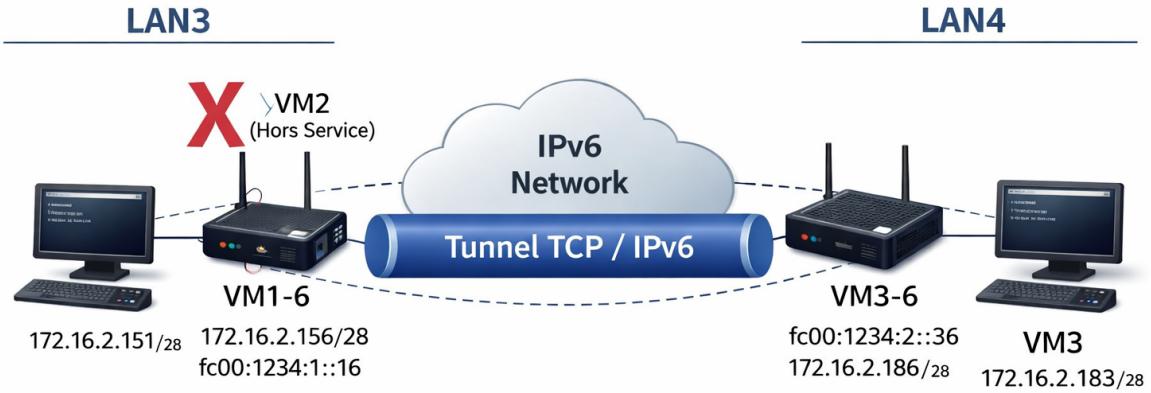


Figure 1: Topology diagram (VM2 down, TCP/IPv6 tunnel between VM1-6 and VM3-6).

```
root@VM1:/vagrant# ping 172.16.2.183
PING 172.16.2.183 (172.16.2.183) 56(84) bytes of data.
```

Figure 2: Ping VM1 → VM3 launched without response while VM2 is down.

```
root@VM1-6:/vagrant/src/extremite# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6b:55:2b brd ff:ff:ff:ff:ff:ff
    altname enp0s3
    inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic eth0
        valid_lft 61735sec preferred_lft 61735sec
    inet6 fe80::a00:27ff:fe6b:552b/64 scope link
        valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:ce:58:67 brd ff:ff:ff:ff:ff:ff
    altname enp0s8
    inet 172.16.2.156/28 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fe80::a00:27ff:fece:5867/64 scope link
        valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:cc:3a:22 brd ff:ff:ff:ff:ff:ff
    altname enp0s9
    inet6 fc00:1234:1::16/64 scope global
        valid_lft forever preferred_lft forever
6: tun0: <POINTOPOINT,MULTICAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 500
    link/none
    inet 172.16.2.1/24 scope global tun0
        valid_lft forever preferred_lft forever
root@VM1-6:/vagrant/src/extremite# ip route
default via 10.0.2.2 dev eth0
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.15
172.16.2.0/24 dev tun0 proto kernel scope link src 172.16.2.1
172.16.2.144/28 dev eth1 proto kernel scope link src 172.16.2.156
172.16.2.176/28 dev tun0 scope link
```

Figure 3: Creation of the tun0 interface.

```

#!/usr/bin/env bash
set -euo pipefail

sudo -s <<'EOS'
modprobe tun 2>/dev/null || true
ip tuntap add dev tun0 mode tun 2>/dev/null || true
ip link set tun0 up

# Address requested by the instructions (example of local config)
ip addr flush dev tun0 || true
ip addr add 172.16.2.1/24 dev tun0
EOS

```

Figure 4 shows the tun0 interface created and active after executing the script.

```

root@VM1-6:/vagrant/src/extremite# ip link show tun0
6: tun0: <POINTOPOINT,MULTICAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 500
    link/none

```

Figure 4: tun0 interface created and active (ip link show tun0 output).

### 3.2.2 Routing: VM1 or VM1-6? (2.2.2)

After the VM2 failure, we must **modify the routing on VM1 and on VM1-6**:

- **VM1**: route to VM3's network (LAN4) via VM1-6 (LAN3).
- **VM1-6**: route from LAN4 to tun0 in order to inject traffic into the tunnel.

This ensures that traffic leaving VM1 is redirected to the endpoint (VM1-6) and then injected into tun0.<sup>5</sup>

### 3.2.3 Local ping tun0 (2.2.3)

Local pings to tun0 addresses are visible in Figures 5 and 6.

```

root@VM1-6:/vagrant/src/extremite# ping 172.16.2.1
PING 172.16.2.1 (172.16.2.1) 56(84) bytes of data.
64 bytes from 172.16.2.1: icmp_seq=1 ttl=64 time=0.013 ms
64 bytes from 172.16.2.1: icmp_seq=2 ttl=64 time=0.032 ms
64 bytes from 172.16.2.1: icmp_seq=3 ttl=64 time=0.036 ms
64 bytes from 172.16.2.1: icmp_seq=4 ttl=64 time=0.033 ms
64 bytes from 172.16.2.1: icmp_seq=5 ttl=64 time=0.033 ms
^C
--- 172.16.2.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4085ms
rtt min/avg/max/mdev = 0.013/0.029/0.036/0.008 ms

```

Figure 5: VM1-6: local ping to 172.16.2.1 (tun0 address).

### 3.2.4 Non-local ping via tun0 (2.2.4) and explanation (2.2.5)

The ping to 172.16.2.10 routed to tun0 without a user-space process is illustrated in Figure 7.

---

<sup>5</sup>Question "2.2.2".

```
Terminal - root@VM1: /vagrant
Fichier Edition Affichage Terminal Onglets Aide
[sudo] Mot de passe de debian-tp :
root@VM1:/home/debian-tp#
root@VM1:/home/debian-tp# ls
Desktop
root@VM1:/home/debian-tp# cd ..
root@VM1:/home# cd ..
root@VM1:# ls
bin etc lib mnt run tmp vmlinuz
boot home lib64 opt sbin usr vmlinuz.old
box.info initrd.img lost+found proc srv vagrant
dev media root sys var
root@VM1:# cd vagrant/
root@VM1:/vagrant# ls
config docs Makefile script src tests VM VM-6
root@VM1:/vagrant# sudo ./tests/test_ipfw tun0 | hexdump -C
[OK] created tun0 (fd=3)
Configure it in another terminal, then ping 172.16.2.10.
This program dumps packets to stdout. Use: | hexdump -C
00000000 45 00 00 54 3d da 40 00 40 01 a8 a3 ac 10 02 01 |E..T=@.0.....
00000010 ac 10 02 0a 08 00 4c e5 aa 00 c9 f9 5b 69 |.....N.....[i
00000020 00 00 00 18 a0 01 b8 00 00 00 00 10 11 12 13 |.....I#.
00000030 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 |.....!#
00000040 24 25 26 27 28 29 2a 2b 2c 2d 2e 2f 30 31 32 33 |$&(!)*,-./0123
root@VM1:#

Terminal - debian-tp@VM1: ~
Fichier Edition Affichage Terminal Onglets Aide
p default qlen 1000
    link/ether 08:00:27:b1:cc:34 brd ff:ff:ff:ff:ff:ff
    altname enp0s8
    inet 172.16.2.131/28 scope global eth1
        valid_lft forever preferred_lft forever
4: eth2 <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group
p default qlen 1000
    link/ether 08:00:27:b6:57:1b brd ff:ff:ff:ff:ff:ff
    altname enp0s9
    inet 172.16.2.151/28 scope global eth2
        valid_lft forever preferred_lft forever
6: tun0 <POINTTOPPOINT,MULTICAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc fq_codel state
UP group default qlen 500
    link/none
    inet 172.16.2.1/28 scope global tun0
        valid_lft forever preferred_lft forever
debian-tp@VM1: ~ sudo ip route replace 172.16.2.10/32 dev tun0
debian-tp@VM1: ~ ping -c 1 172.16.2.10
PING 172.16.2.10 (172.16.2.10) 56(84) bytes of data.

-- 172.16.2.10 ping statistics --
1 packets transmitted, 0 received, 100% packet loss, time 0ms

debian-tp@VM1: ~
```

Figure 6: VM3-6: local ping to 172.16.2.10 (tun0 address).

Figure 7: Left: `test_iftun tun0 | hexdump -C`. Right: ping 172.16.2.10 routed to `tun0` with 100% loss (no response).

### 3.3 Packet Retrieval (2.3)

### 3.3.1 Option IFF\_NO\_PI (2.3.4)

**Role:** `IFF_NO_PI` asks the kernel not to prefix packets read on `tun0` with a PI header. **Without `IFF_NO_PI`:** readings begin with extra PI bytes; a naive IPv4 parser would see its data shifted.<sup>6</sup>

## 4 Simple IPv4 Tunnel over TCP/IPv6 (Part 3)

#### 4.1 Incoming Traffic Redirection (3.1)

The test of the `ext-out` server with TCP/IPv6 reception redirected to `stdout` is visible in Figure 8.

## 4.2 Final Integration (3.3)

The setup of the bidirectional tunnel is visible in Figure 9.

### 4.3 VM1-6 <-> VM3-6 Tunnel: Diagrams (3.4)

The complete diagram is provided in Figure 10.

## 5 Functional Validation (Part 4)

## 5.1 Configuration (4.1)

On VM1-6, VM1, VM2-6, VM3-6, and VM3: display ip addr and ip route.<sup>7</sup>

<sup>6</sup>Question "2.3.4".

<sup>7</sup>Point "4.1".

```

^CStopping tunnel...
root@VM1-6:/vagrant/src/extremite# sudo python3 test_main.py server
Interface tun0 created successfully.

-----
Interface created. RUN ./configure-tun.sh NOW!
Press ENTER once configured.

-----
Server listening on port 123...
Connection accepted from ('fc00:1234:2::36', 51876, 0, 0)
--- Tunnel actif ---
--> TUN: Lu 76 octets. Envoi...
<- NET: Reçu 76 octets. Ecriture...
--> TUN: Lu 76 octets. Envoi...
<- NET: Reçu 2048 octets. Ecriture...
<- NET: Reçu 808 octets. Ecriture...
ERREUR CRITIQUE: [Errno 22] Invalid argument
root@VM1-6:/vagrant/src/extremite#

```

Figure 8: ext-out test: reception via TCP/IPv6 redirected to stdout.

```

root@VM3-6:/vagrant/src/extremite# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host
            valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6b:55:2b brd ff:ff:ff:ff:ff:ff
        altname enp0s3
        inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic eth0
            valid_lft 61597sec preferred_lft 61597sec
        inet6 fe80::a00:27ff:fe6b:552b/64 scope link
            valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:0a:d5:3b brd ff:ff:ff:ff:ff:ff
        altname enp0s8
        inet 172.16.2.186/28 scope global eth1
            valid_lft forever preferred_lft forever
        inet6 fe80::a00:27ff:fe0a:d53b/64 scope link
            valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:76:8d:56 brd ff:ff:ff:ff:ff:ff
        altname enp0s9
        inet6 fc00:1234:2::36/64 scope global
            valid_lft forever preferred_lft forever
7: tun0: <POINTOPOINT,MULTICAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 500
    link/none
    inet 172.16.2.10/24 scope global tun0
        valid_lft forever preferred_lft forever
root@VM3-6:/vagrant/src/extremite# ip route
default via 10.0.2.2 dev eth0
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.15
172.16.2.0/24 dev tun0 proto kernel scope link src 172.16.2.10
172.16.2.144/28 dev tun0 scope link
172.16.2.176/28 dev eth1 proto kernel scope link src 172.16.2.186
root@VM3-6:/vagrant/src/extremite#

```

Figure 9: TCP/IPv6 connection established and bidirectional tunnel active (test\_main.py).

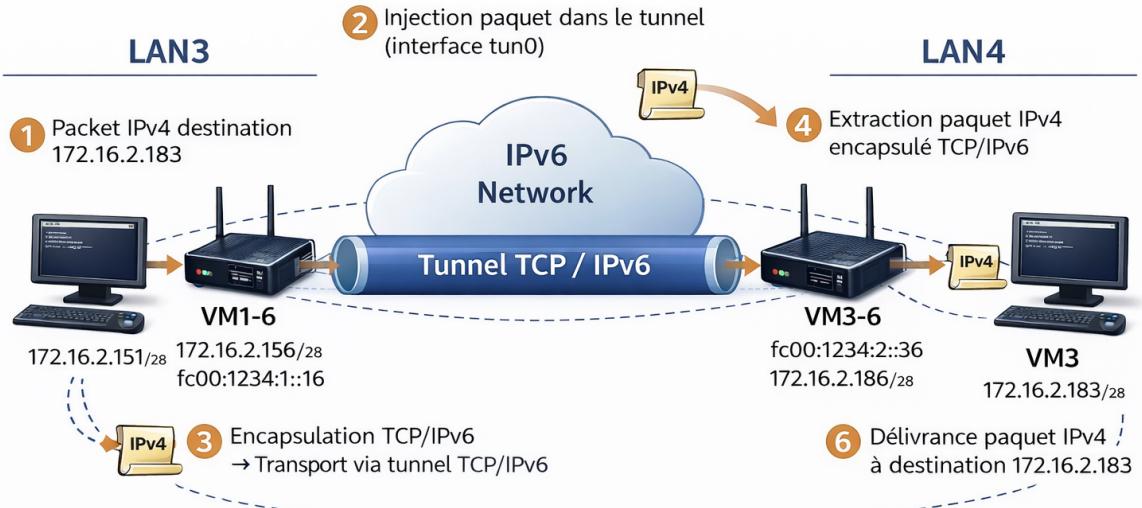


Figure 10: Complete path of an IPv4 packet from VM1 to VM3 via the TCP/IPv6 tunnel.

```
root@VM1-6:/vagrant/src/extremite# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6b:55:2b brd ff:ff:ff:ff:ff:ff
    altname enp0s3
    inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic eth0
        valid_lft 61735sec preferred_lft 61735sec
    inet6 fe80::a00:27ff:fe6b:552b/64 scope link
        valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:ce:58:67 brd ff:ff:ff:ff:ff:ff
    altname enp0s8
    inet 172.16.2.156/28 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fe80::a00:27ff:fece:5867/64 scope link
        valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:cc:3a:22 brd ff:ff:ff:ff:ff:ff
    altname enp0s9
    inet6 fc00:1234:1::1/64 scope global
        valid_lft forever preferred_lft forever
6: tun0: <POINTOPOINT,MULTICAST,NOARP,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 500
    link/none
    inet 172.16.2.1/24 scope global tun0
        valid_lft forever preferred_lft forever
root@VM1-6:/vagrant/src/extremite# ip route
default via 10.0.2.2 dev eth0
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.15
172.16.2.0/24 dev tun0 proto kernel scope link src 172.16.2.1
172.16.2.144/28 dev eth1 proto kernel scope link src 172.16.2.156
172.16.2.176/28 dev tun0 scope link
```

Figure 11: VM1-6 ip addr and ip route.

```
vagrant@VM2-6:~$ ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host
            valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6b:55:2b brd ff:ff:ff:ff:ff:ff
        altname enp0s3
        inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic eth0
            valid_lft 61674sec preferred_lft 61674sec
        inet6 fe80::a00:27ff:fe6b:552b/64 scope link
            valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:7d:40:9a brd ff:ff:ff:ff:ff:ff
        altname enp0s8
        inet6 fc00:1234:1::26/64 scope global
            valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:a7:d0:d4 brd ff:ff:ff:ff:ff:ff
        altname enp0s9
        inet6 fc00:1234:2::26/64 scope global
            valid_lft forever preferred_lft forever
vagrant@VM2-6:~$ ip route
default via 10.0.2.2 dev eth0
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.15
vagrant@VM2-6:~$
```

Figure 12: VM2-6 (IPv4 router) `ip addr` and `ip route`.

```
root@VM1:/vagrant# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6b:55:2b brd ff:ff:ff:ff:ff:ff
        altname enp0s3
        inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic eth0
            valid_lft 61664sec preferred_lft 61664sec
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:b1:cc:34 brd ff:ff:ff:ff:ff:ff
        altname enp0s8
        inet 172.16.2.131/28 scope global eth1
            valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:b6:57:1b brd ff:ff:ff:ff:ff:ff
        altname enp0s9
        inet 172.16.2.151/28 scope global eth2
            valid_lft forever preferred_lft forever
root@VM1:/vagrant# ip route
default via 10.0.2.2 dev eth0
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.15
172.16.2.128/28 dev eth1 proto kernel scope link src 172.16.2.131
172.16.2.144/28 dev eth2 proto kernel scope link src 172.16.2.151
172.16.2.160/28 via 172.16.2.132 dev eth1
172.16.2.176/28 via 172.16.2.156 dev eth2
root@VM1:/vagrant#
```

Figure 13: VM1 `ip addr` and `ip route`.

```
root@VM3:/vagrant/script/net# python3 -u - <<'PY'
import socket
s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.bind(("0.0.0.0",12345))
s.listen(1)
c,a=s.accept()
print("client:",a, flush=True)
data=c.recv(4096)
print("RECU:", data.decode(errors="replace"), flush=True)
PY
client: ('172.16.2.151', 47320)
RECU: HELLO_TUNNEL

root@VM3:/vagrant/script/net# python3 -u - <<'PY'
import socket
s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.bind(("0.0.0.0",12345))
s.listen(1)
c,a=s.accept()
print("client:",a, flush=True)
data=c.recv(4096)
print("RECU:", data.decode(errors="replace"), flush=True)
PY

client: ('172.16.2.151', 47880)
RECU: HELLO_TUNNEL
```

Figure 14: VM3 ip addr and ip route.

## 5.2 Layer 3 (4.2)

From VM1: ping 172.16.2.183.<sup>8</sup>

```
root@VM3-6:/vagrant/src/extremite# ping 172.16.2.10
PING 172.16.2.10 (172.16.2.10) 56(84) bytes of data.
64 bytes from 172.16.2.10: icmp_seq=1 ttl=64 time=0.014 ms
64 bytes from 172.16.2.10: icmp_seq=2 ttl=64 time=0.033 ms
64 bytes from 172.16.2.10: icmp_seq=3 ttl=64 time=0.030 ms
64 bytes from 172.16.2.10: icmp_seq=4 ttl=64 time=0.034 ms
^C
--- 172.16.2.10 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3066ms
rtt min/avg/max/mdev = 0.014/0.027/0.034/0.008 ms
```

Figure 15: Ping VM1 → VM3 (172.16.2.183) and ICMP capture on tun0 side VM1-6/VM3-6.

## 5.3 Layer 4 (4.3)

From VM1: create msg.txt and send it to VM3 (TCP service).<sup>9</sup>

```
root@VM3:/vagrant/script/net# python3 -u - <<'PY'
import socket
s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.bind(("0.0.0.0",12345))
s.listen(1)
c,a=s.accept()
print("client:",a, flush=True)
data=c.recv(4096)
print("RECU:", data.decode(errors="replace"), flush=True)
PY
client: ('172.16.2.151', 47320)
RECU: HELLO_TUNNEL

root@VM3:/vagrant/script/net# python3 -u - <<'PY'
import socket
s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.bind(("0.0.0.0",12345))
s.listen(1)
c,a=s.accept()
print("client:",a, flush=True)
data=c.recv(4096)
print("RECU:", data.decode(errors="replace"), flush=True)
PY

client: ('172.16.2.151', 47880)
RECU: HELLO_TUNNEL
```

Figure 16: VM1: creation of msg.txt and sending to VM3 (TCP client).

<sup>8</sup>Point "4.2".

<sup>9</sup>Point "4.3".

## 5.4 Layer 4: Bandwidth (4.4)

Iperf3 server on VM3 then client on VM1-6 with multiple buffer sizes: 10, 2K, 128K, 1M.<sup>10</sup>

Figure 17: Iperf3 results buffers ( 10 ): measured throughput.

Figure 18: Iperf3 results buffers ( 1M): measured throughput.

## 6 Conclusion

The IPv4 tunnel encapsulated within TCP/IPv6 restores connectivity between LAN3 and LAN4 despite the VM2 failure. The required tests (configuration, ping, application TCP, iperf3) are satisfied and the captures confirm the effective circulation of packets via `tun0` and TCP/IPv6 transport.

<sup>10</sup>Point "4.4".